

# Operating Systems

## Lecture 3

Computer Security DD2395

Roberto Guanciale

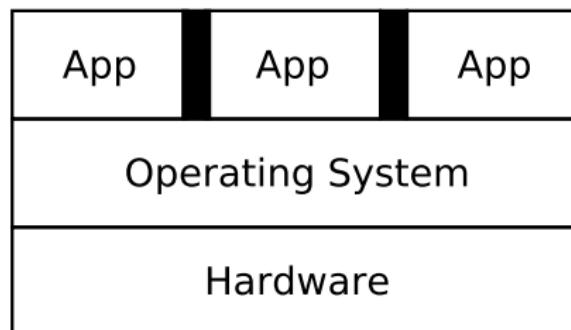
[robertog@kth.se](mailto:robertog@kth.se)

2017-11-03

- Assistant professor Computer Security
- Research interest in Formal Verification
  - Hack-proof Operating Systems
- Enthusiast software developer
- robertog@kth.se (use [DD2395] in mail subject)
- Office: Building D, level 5, room 4529
- You are welcomed (no booking required, no fee)
  - Tuesday 13:15-14:30

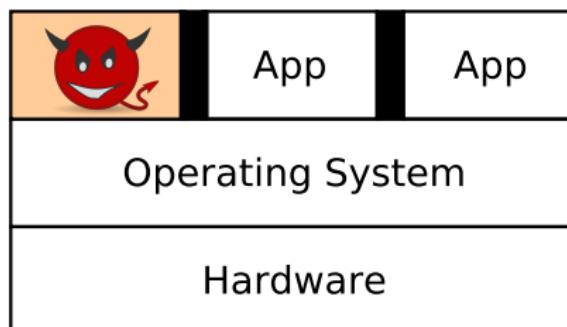
# Operating Systems

- An OS is a program that acts an intermediary between applications and computer hardware
- Makes application software portable and versatile.
- Allows sharing of hardware and software resources.
- Provides isolation, security and protection among applications.



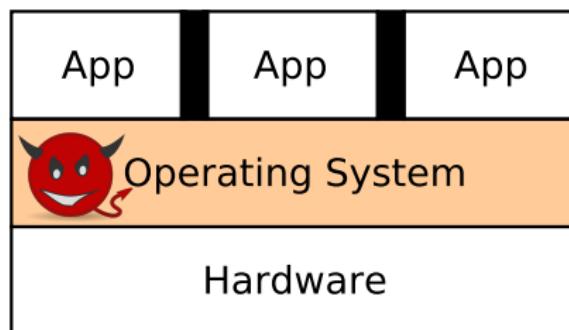
# Operating Systems

- An OS is a program that acts an intermediary between applications and computer hardware
- Makes application software portable and versatile.
- Allows sharing of hardware and software resources.
- Provides isolation, security and protection among applications.



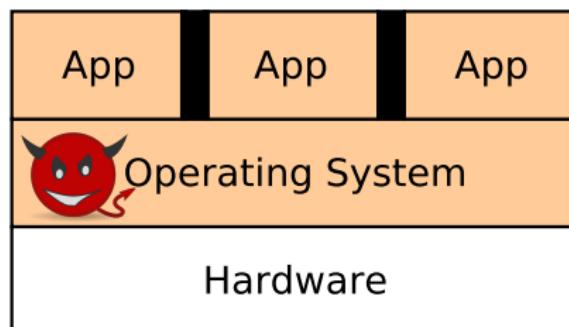
# Operating Systems

- An OS is a program that acts an intermediary between applications and computer hardware
- Makes application software portable and versatile.
- Allows sharing of hardware and software resources.
- Provides isolation, security and protection among applications.



# Operating Systems

- An OS is a program that acts an intermediary between applications and computer hardware
- Makes application software portable and versatile.
- Allows sharing of hardware and software resources.
- Provides isolation, security and protection among applications.



# Main questions

- How can we run Chrome on a modern architecture?
- How can we isolate Chrome from LibreOffice?
- How can we run multiple instances of Chrome?
- How can we interact with peripherals?

## Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)

## Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)
- S40 (~ 1 millions LOC)

## Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)
- S40 (~ 1 millions LOC)
- Free RTOS (~ 10 000 LOC)

## Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)
- S40 (~ 1 millions LOC)
- Free RTOS (~ 10 000 LOC)



# Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)
- S40 (~ 1 millions LOC)
- Free RTOS (~ 10 000 LOC)



# Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)
- S40 (~ 1 millions LOC)
- Free RTOS (~ 10 000 LOC)



# Example of OSes

- Windows/Linux/OS X (~ 50 millions LOC)
- Android/iOS (~ 50 millions LOC)
- S40 (~ 1 millions LOC)
- Free RTOS (~ 10 000 LOC)



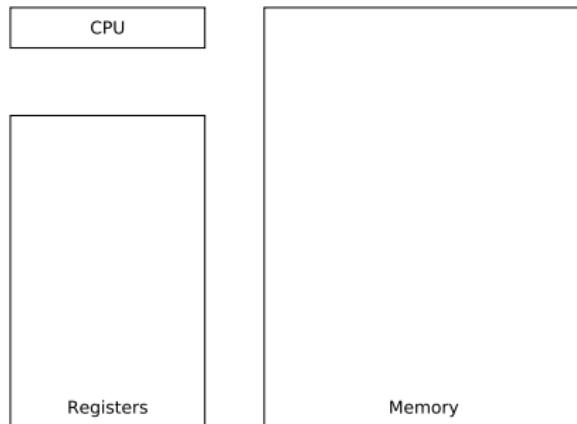
# OS by example

How can we run Chrome on a modern CPU?



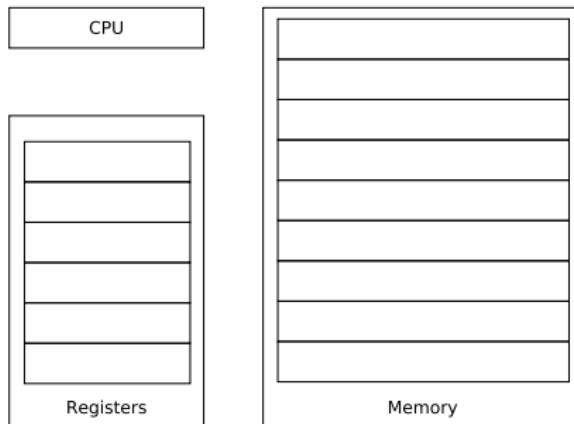
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC



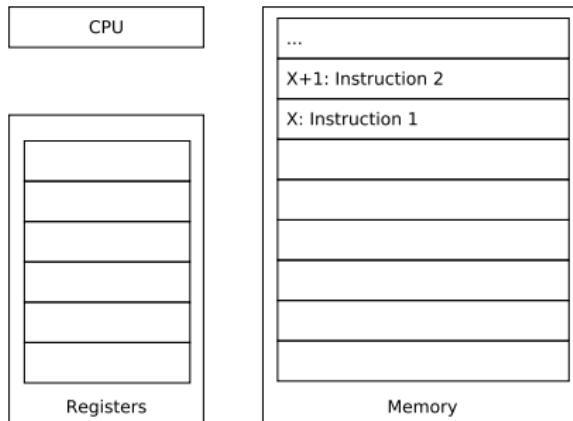
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC



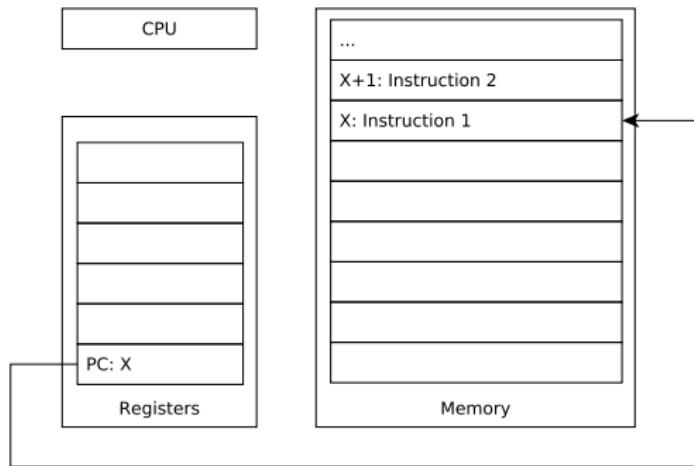
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory



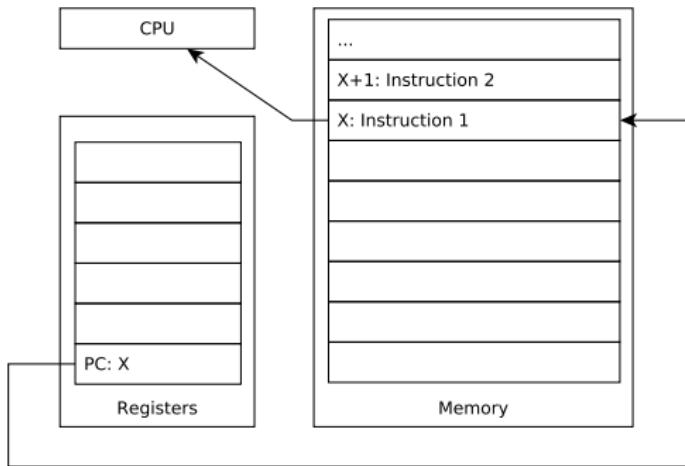
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory
- Program Counter (PC): address of the current instruction



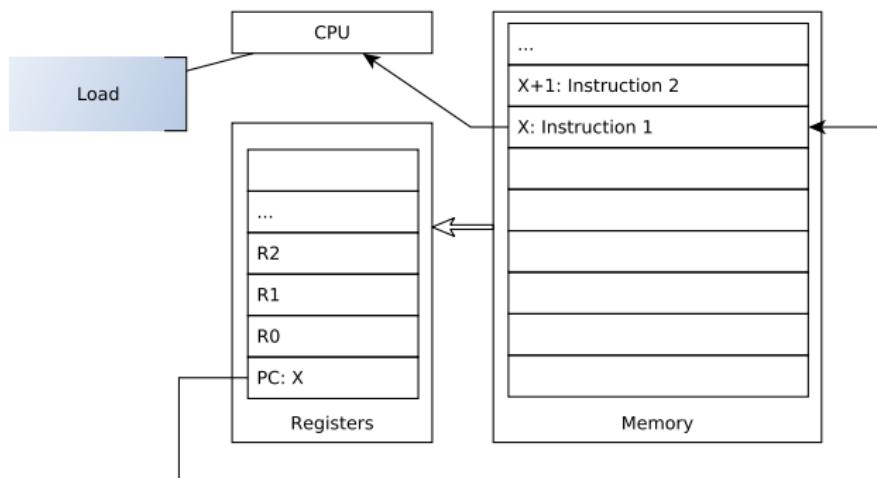
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory
- Program Counter (PC): address of the current instruction



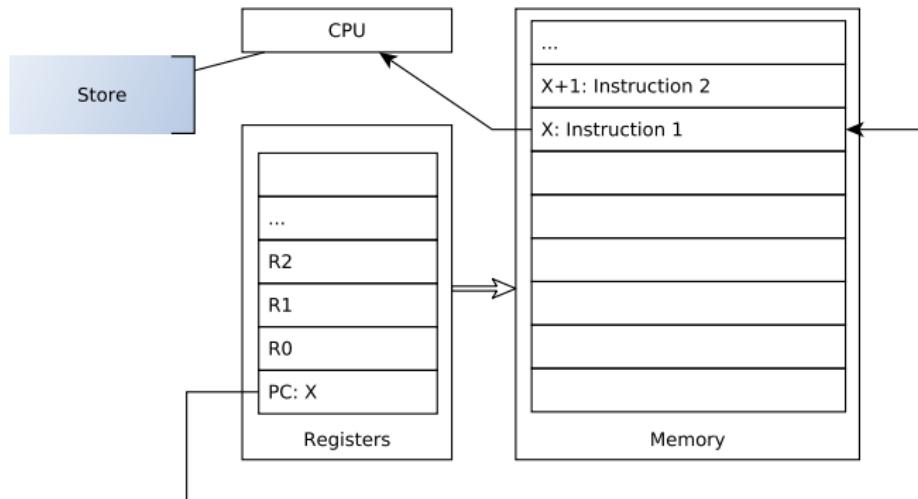
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory
- Program Counter (PC): address of the current instruction



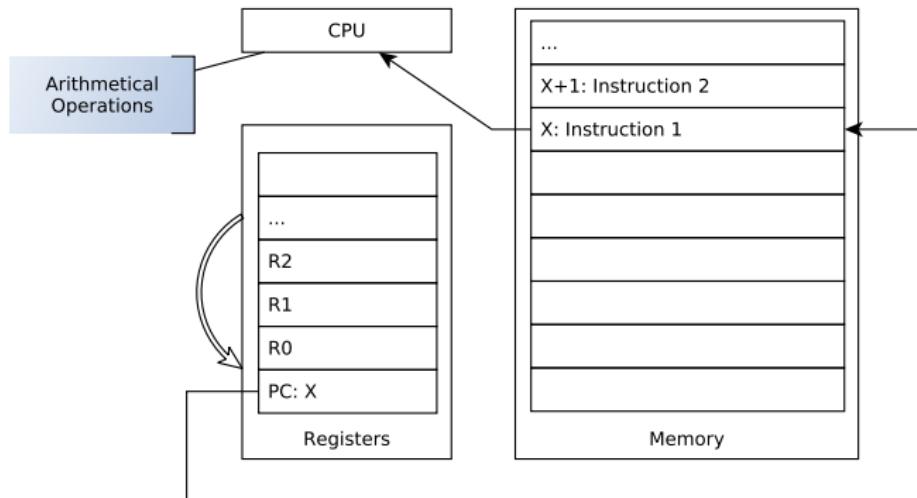
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory
- Program Counter (PC): address of the current instruction



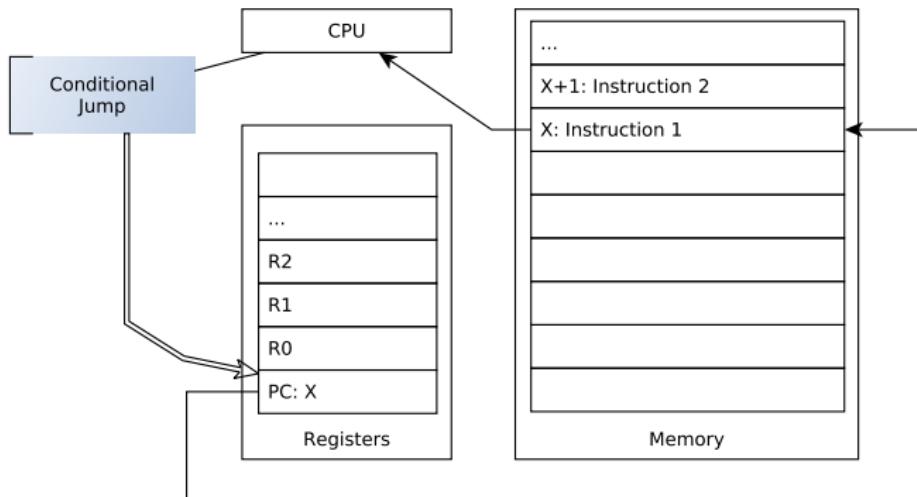
# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory
- Program Counter (PC): address of the current instruction



# HW architecture

- Common architectures x86, x64, ARM, MIPS, PowerPC
- Program stored in memory
- Program Counter (PC): address of the current instruction



# Application memory

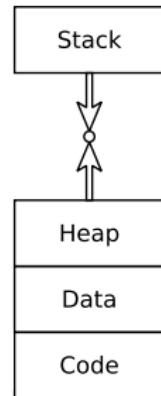
```
int x;
int * hello(int y) {
    int * res = malloc(8);
    res[0] = x + 10;
    res[1] = y + x;
    return res;
}
```

# Application memory

```
class Example {  
    int x;  
    static int x;  
    int * hello(int y) {  
        int * res = malloc(8);  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
}  
  
int[] hello(int y) {  
    int[] res = new int[2];  
    res[0] = x + 10;  
    res[1] = y + x;  
    return res;  
}
```

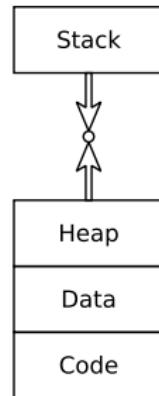
# Application Memory

```
class Example {  
    int x;  
    int * hello(int y) {  
        int * res = malloc(8);  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
};  
  
class Example {  
    static int x;  
    int[] hello(int y) {  
        int[] res = new int[2];  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
};
```



# Application Memory

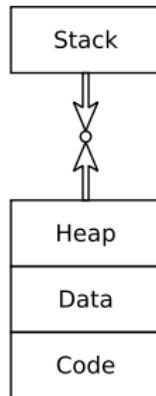
```
class Example {  
    int x;  
    int * hello(int y) {  
        int * res = malloc(8);  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
    static int x;  
    int[] hello(int y) {  
        int[] res = new int[2];  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
}
```



Discuss five minutes with your colleagues  
Where each element is allocated?

# Application Memory

```
class Example {  
    int x;  
    int * hello(int y) {  
        int * res = malloc(8);  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
};  
  
class Example {  
    static int x;  
    int[] hello(int y) {  
        int[] res = new int[2];  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
};
```

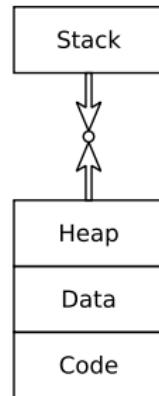


Question 2: Where does malloc allocate memory?

- A - Code
- B - Data
- C - Heap
- D - Stack

# Application Memory

```
class Example {  
    int x;  
    int * hello(int y) {  
        int * res = malloc(8);  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
};  
  
class Example {  
    static int x;  
    int[] hello(int y) {  
        int[] res = new int[2];  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
};
```



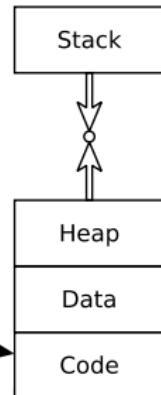
Question 3: Where is y allocated?

- A - Code
- B - Data
- C - Heap
- D - Stack

# Application Memory

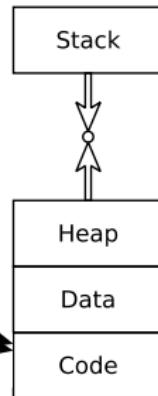
```
int x;  
int * hello(int y) {  
    int * res = malloc(8);  
    res[0] = x + 10;  
    res[1] = y + x;  
    return res;  
}
```

```
class Example {  
    static int x;  
    int[] hello(int y) {  
        int[] res = new int[2];  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
}
```



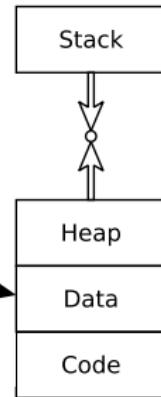
# Application Memory

```
int x;
class Example {
    static int x;
    int * hello(int y) {
        int * res = malloc(8);
        res[0] = x + 10;
        res[1] = y + x;
        return res;
    }
    int[] hello(int y) {
        int[] res = new int[2];
        res[0] = x + 10;
        res[1] = y + x;
        return res;
    }
}
```



# Application Memory

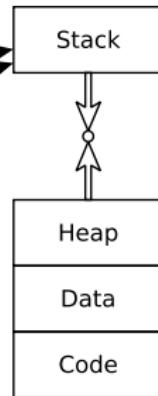
```
int x;  
int * hello(int y) {  
    int * res = malloc(8);  
    res[0] = x + 10;  
    res[1] = y + x;  
    return res;  
}  
  
class Example {  
    static int x;  
    int[] hello(int y) {  
        int[] res = new int[2];  
        res[0] = x + 10;  
        res[1] = y + x;  
        return res;  
    }  
}
```



# Application Memory

```
int x;
int * hello(int y) {
    int * res = malloc(8);
    res[0] = x + 10;
    res[1] = y + x;
    return res;
}

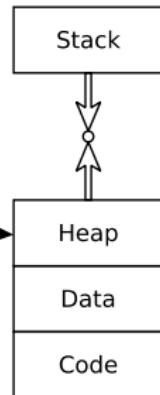
class Example {
    static int x;
    int[] hello(int y) {
        int[] res = new int[2];
        res[0] = x + 10;
        res[1] = y + x;
        return res;
    }
}
```



# Application Memory

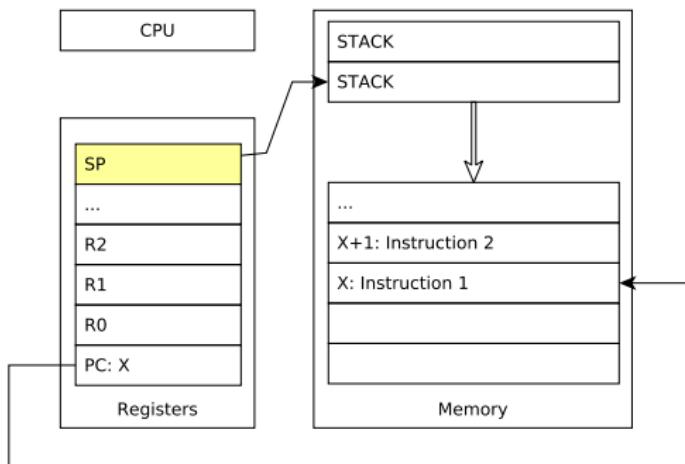
```
int x;
int * hello(int y) {
    int * res = malloc(8);
    res[0] = x + 10;
    res[1] = y + x;
    return res;
}

class Example {
    static int x;
    int[] hello(int y) {
        int[] res = new int[2];
        res[0] = x + 10;
        res[1] = y + x;
        return res;
    }
}
```



# Application Memory

- Stack Pointer points to the bottom of the stack



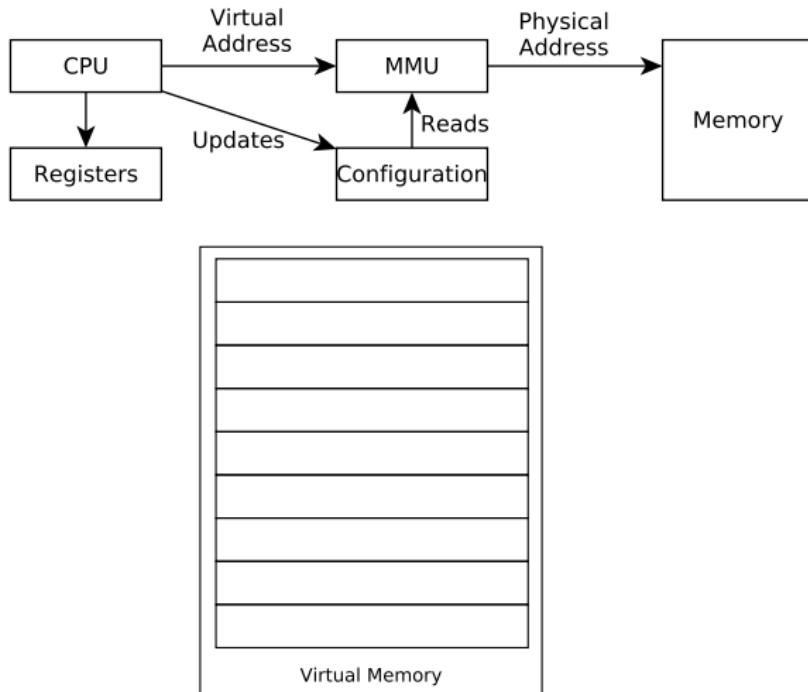
# Virtual Memory

- Modern systems do not use directly physical addresses
  - Run the same software in systems with 1 GB or 2 GB of memory
  - Guarantee that two independent applications do not use the same addresses

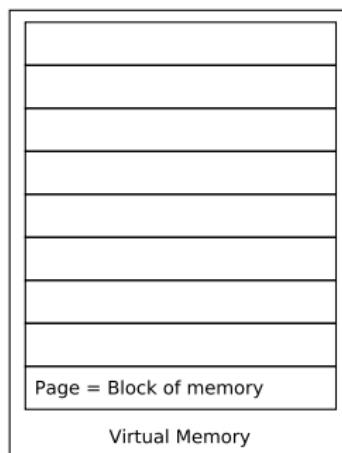
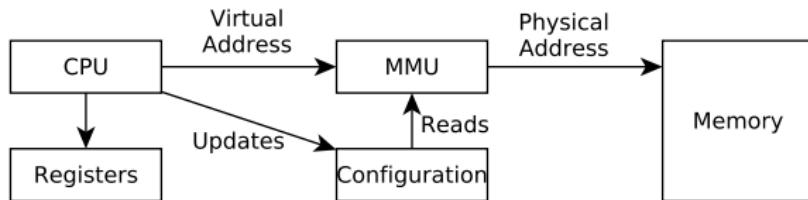
# Virtual Memory

- Modern systems do not use directly physical addresses
  - Run the same software in systems with 1 GB or 2 GB of memory
  - Guarantee that two independent applications do not use the same addresses
- Virtual memory
  - SW uses virtual addresses
  - HW configuration maps virtual addresses to physical addresses
  - HW configuration updatable
  - HW mechanism is called Memory Management Unit

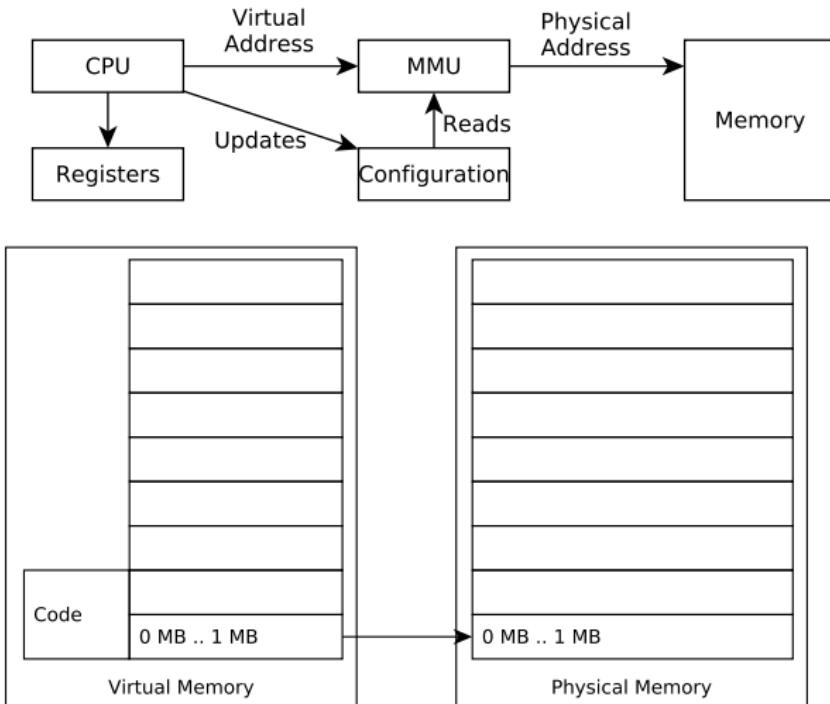
# Virtual Memory



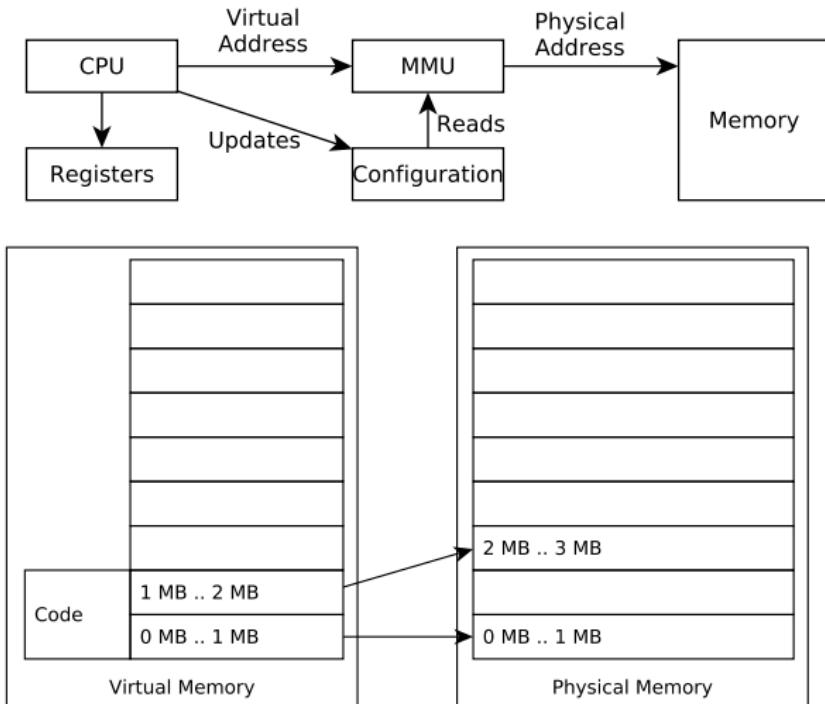
# Virtual Memory



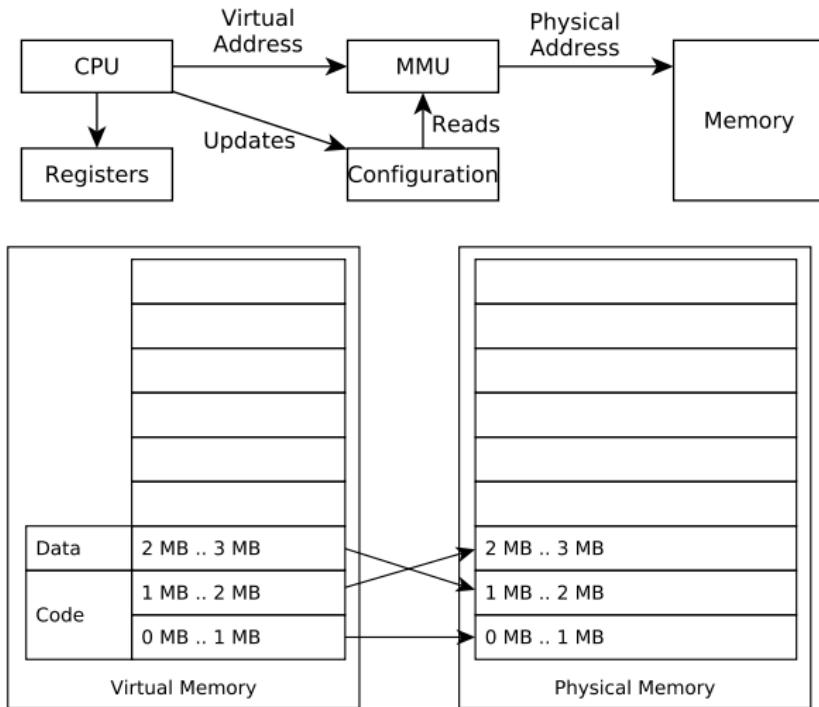
# Virtual Memory



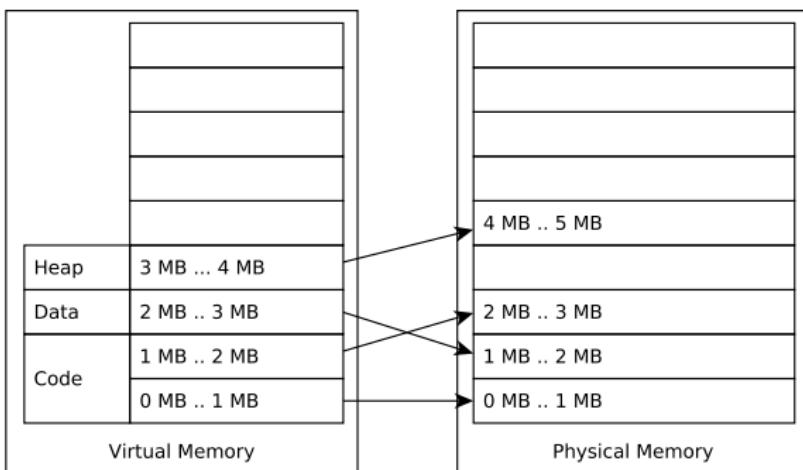
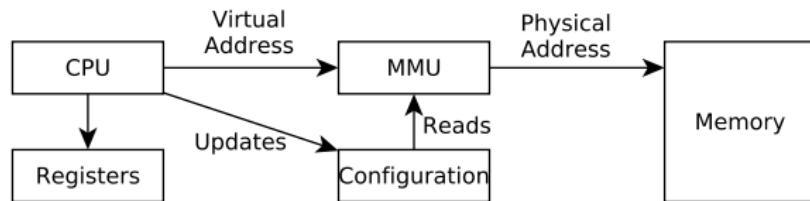
# Virtual Memory



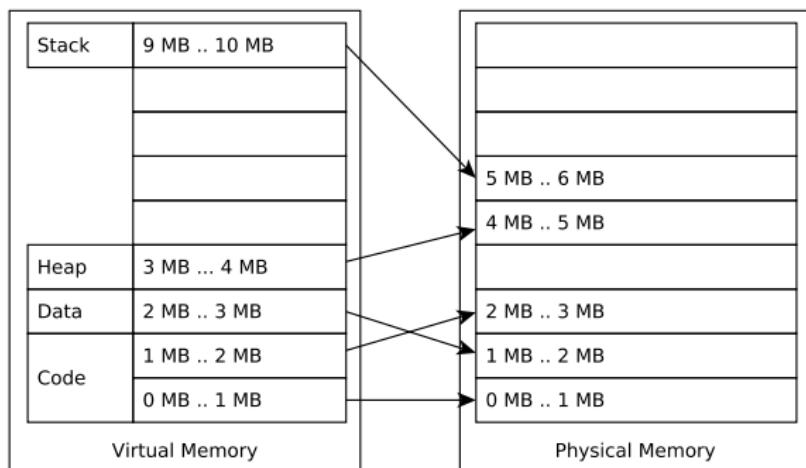
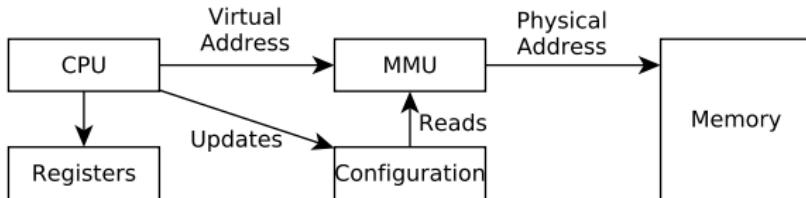
# Virtual Memory



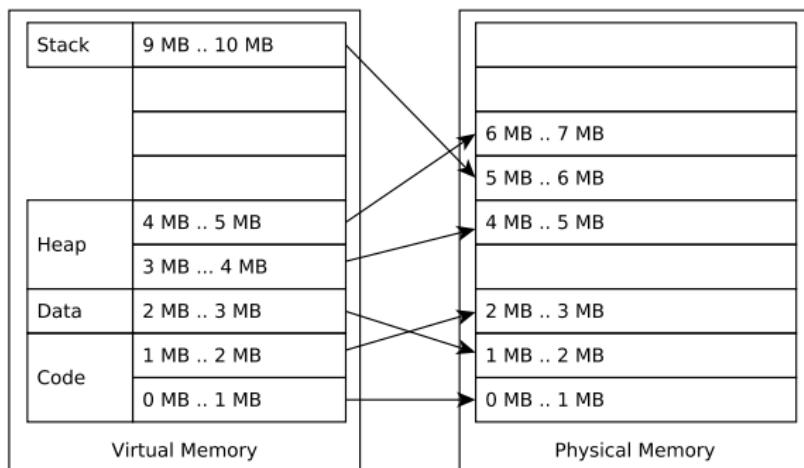
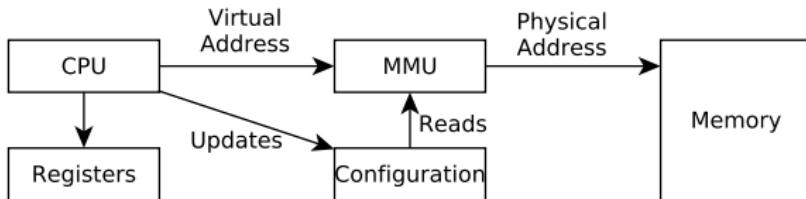
# Virtual Memory



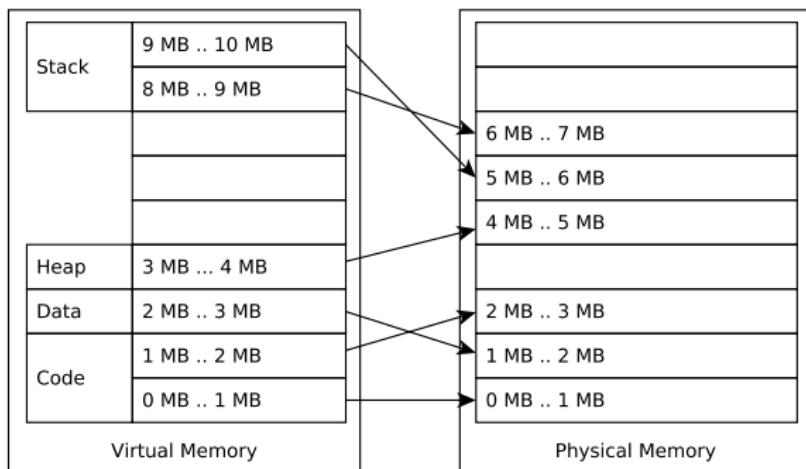
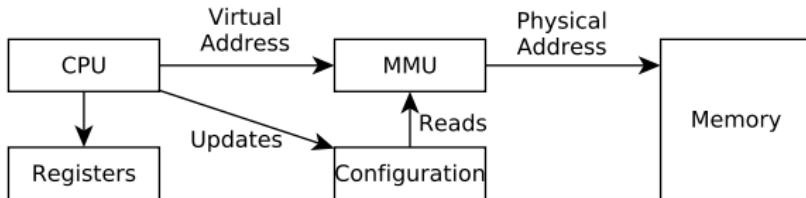
# Virtual Memory



# Virtual Memory



# Virtual Memory



## OS by example: 2

- How can we isolate Chrome from LibreOffice?



## OS by example: 2

- How can we isolate Chrome from LibreOffice?
- Two types of resources:
  - registers
  - memory

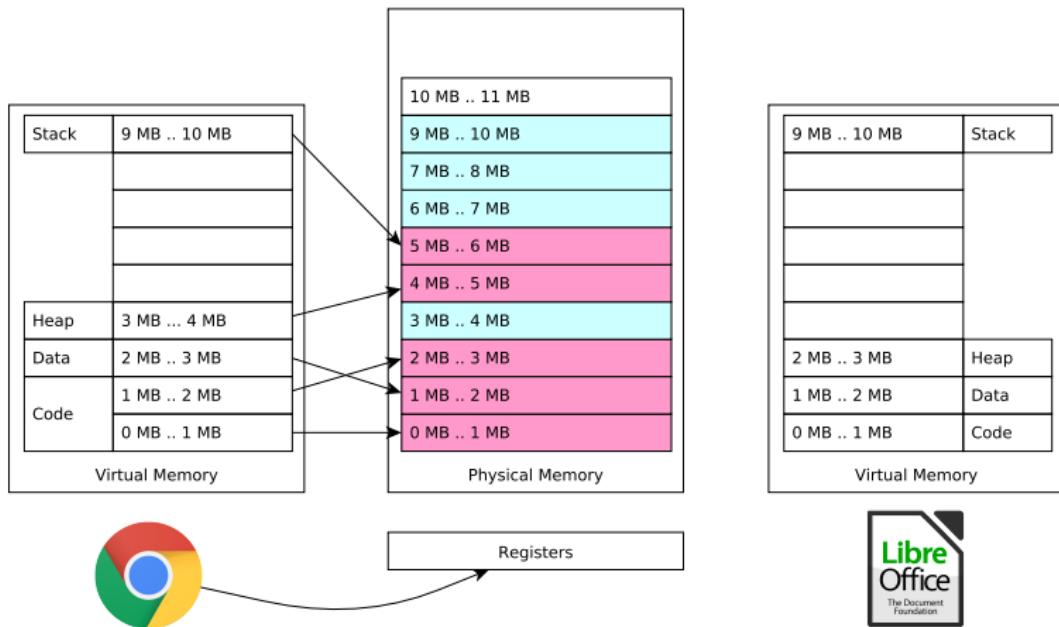


## OS by example: 2

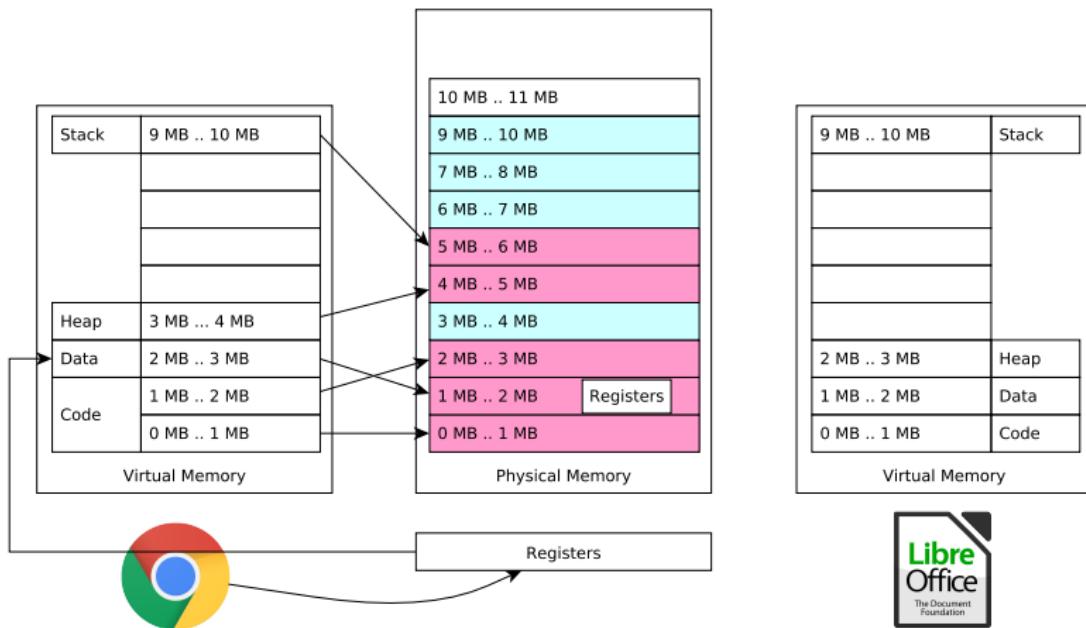
- How can we isolate Chrome from LibreOffice?
- Two types of resources:
  - registers ⇒ multiplex
  - memory ⇒ isolation
- One of the main tasks of an OS



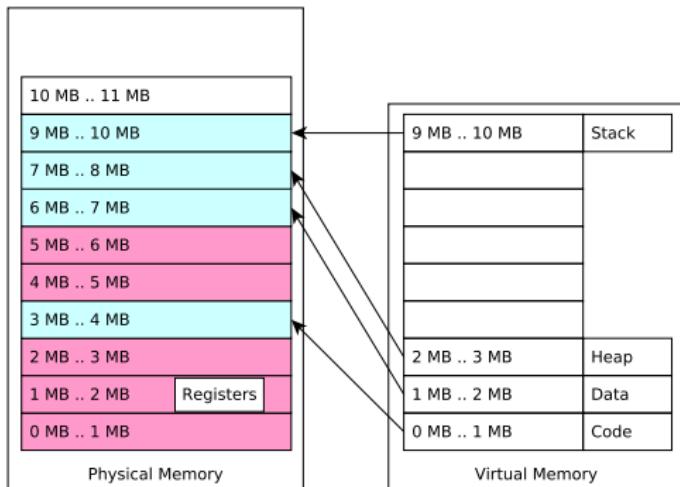
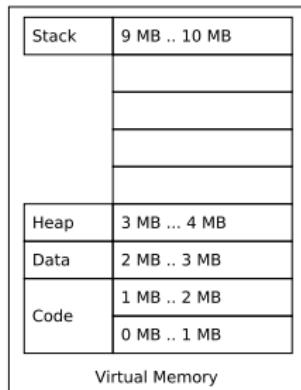
# Context switch



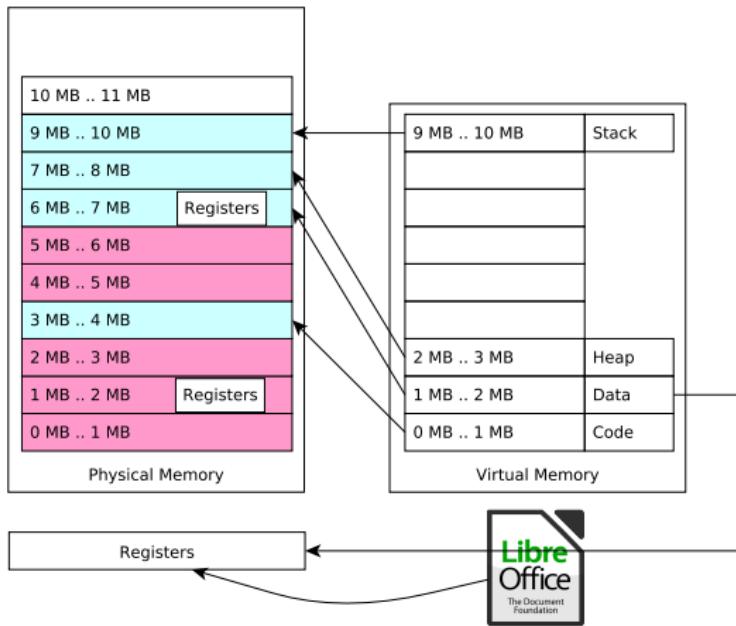
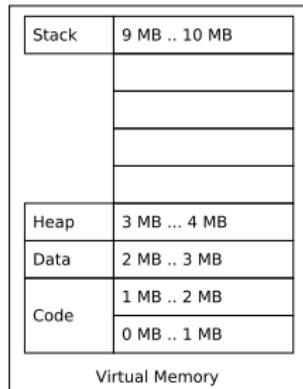
# Context switch



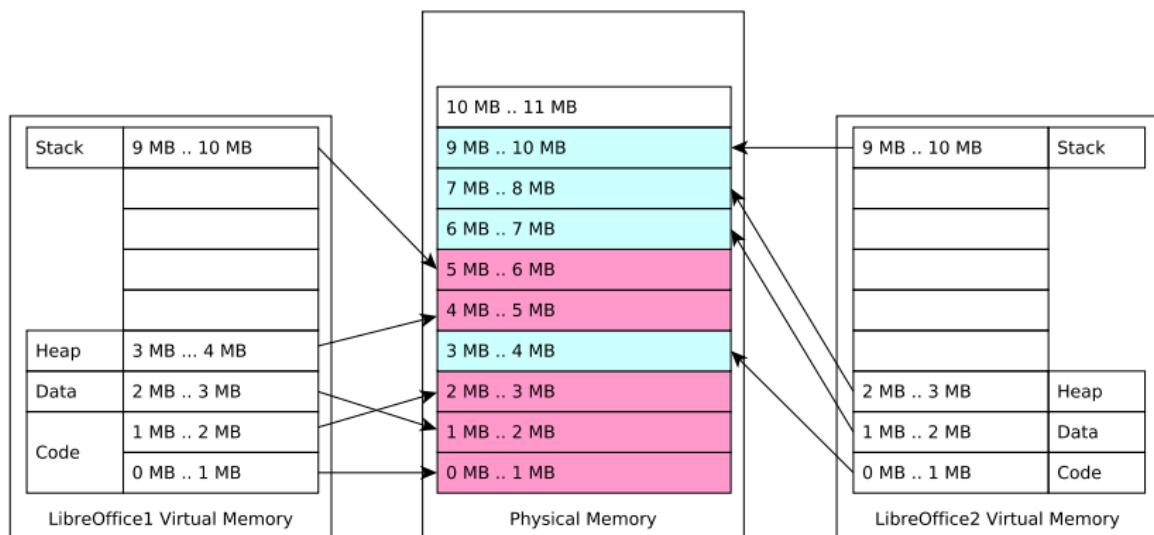
# Context switch



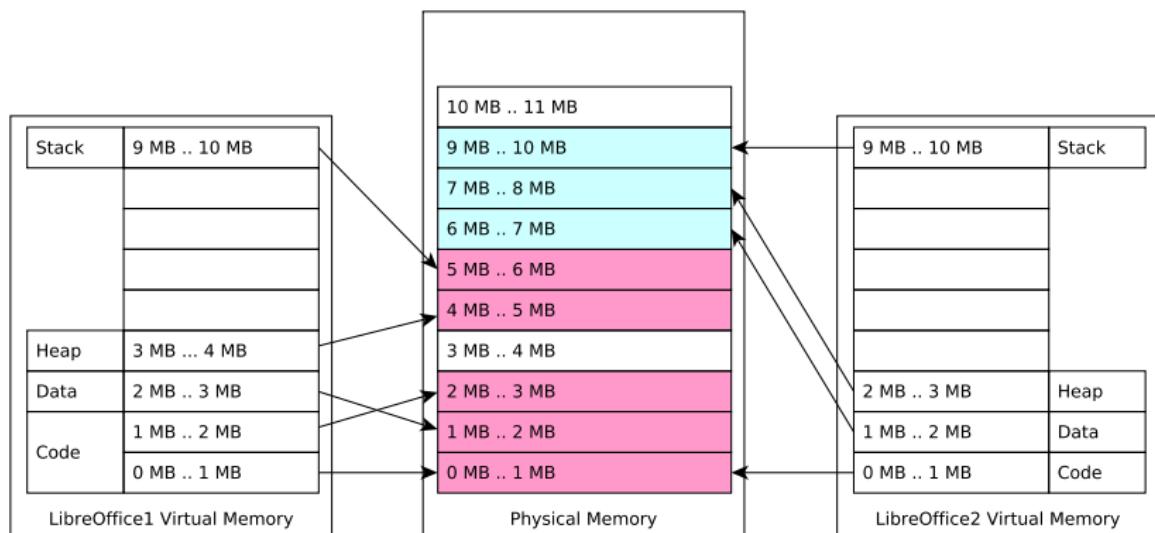
# Context switch



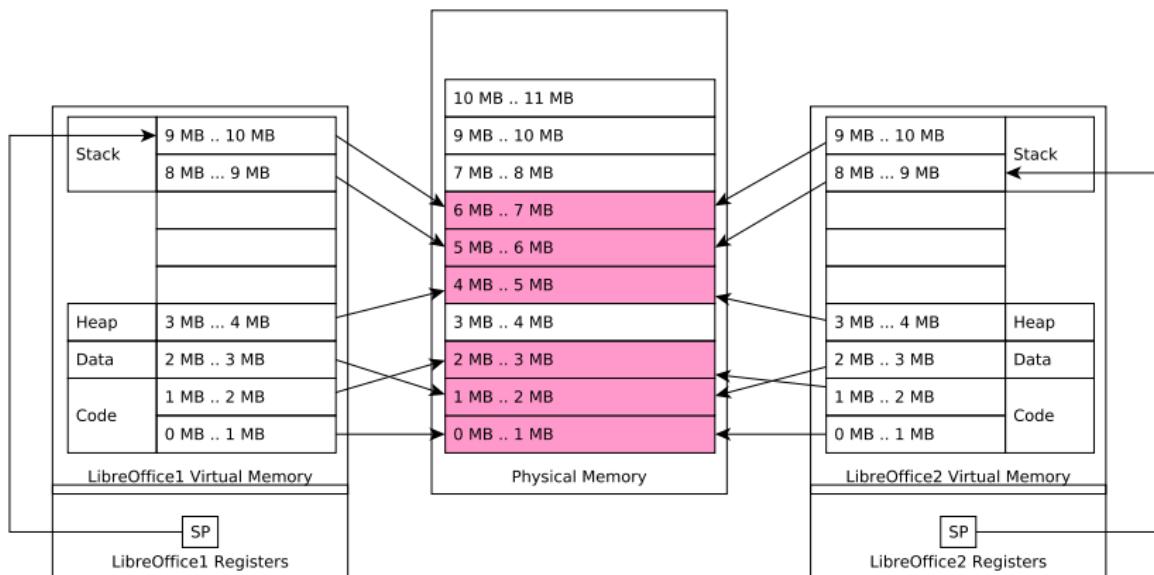
# Processes



# Processes

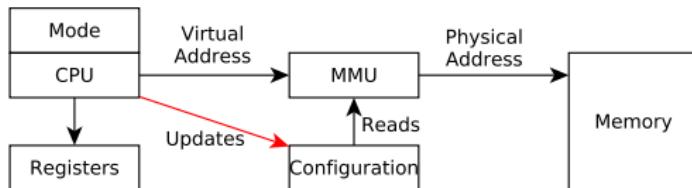


# Threads



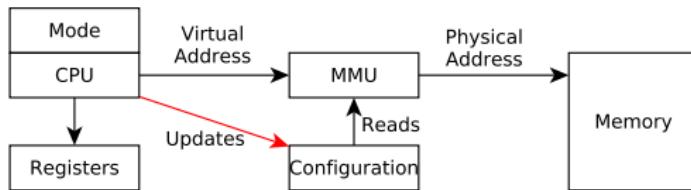
# HW architecture

- Execution mode (i.e. restricted/un-restricted)
- OS takes control of un-restricted mode
- Everything else uses restricted mode



# HW architecture

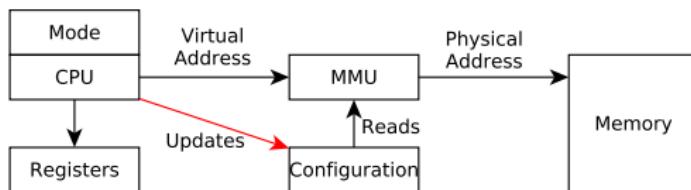
- Execution mode (i.e. restricted/un-restricted)
- OS takes control of un-restricted mode
- Everything else uses restricted mode



- From un-restricted to restricted = exiting the OS
  - instruction to change mode

# HW architecture

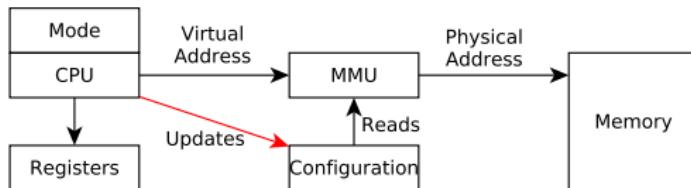
- Execution mode (i.e. restricted/un-restricted)
- OS takes control of un-restricted mode
- Everything else uses restricted mode



- From un-restricted to restricted = exiting the OS
  - instruction to change mode
- From restricted to un-restricted = entering the OS?

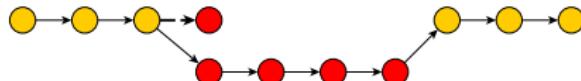
# HW architecture

- Execution mode (i.e. restricted/un-restricted)
- OS takes control of un-restricted mode
- Everything else uses restricted mode



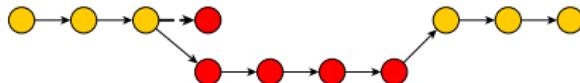
- From un-restricted to restricted = exiting the OS
  - instruction to change mode
- From restricted to un-restricted = entering the OS?
  - interrupts

# SW interrupt



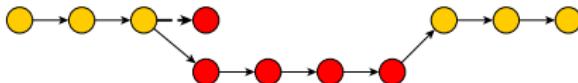
- ① Application executes the Software Interrupt Instruction

# SW interrupt



- ① Application executes the Software Interrupt Instruction
- ② HW
  - Mode switches to un-restricted
  - Program counter jumps to the OS **vector table**

# SW interrupt

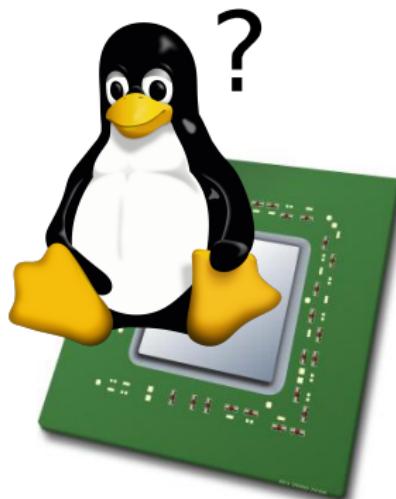


- ① Application executes the Software Interrupt Instruction
- ② HW
  - Mode switches to un-restricted
  - Program counter jumps to the OS **vector table**
- ③ OS
  - Saves the registers (context) in the process memory
  - Jumps to the **exception handler**
  - ...  
(e.g. change MMU configuration and “logical” active process)
  - Loads the registers (context) from the memory of the **active** process
  - Switches the mode to restricted

# OS Memory

① Question 4: How does the OS access the memory?

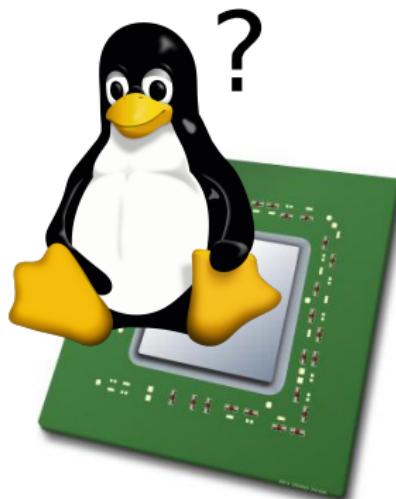
- A - Using directly physical addresses
- B - Using the same virtual addresses of processes
- C - Using dedicated virtual addresses



# OS Memory

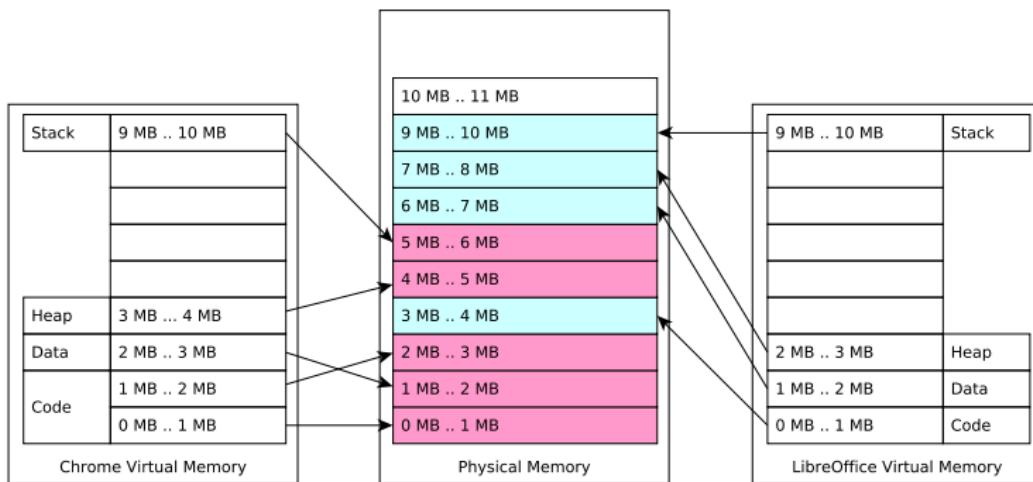
## ① Question 4: How does the OS access the memory?

- A - Using directly physical addresses
- B - Using the same virtual addresses of processes
- C - Using dedicated virtual addresses



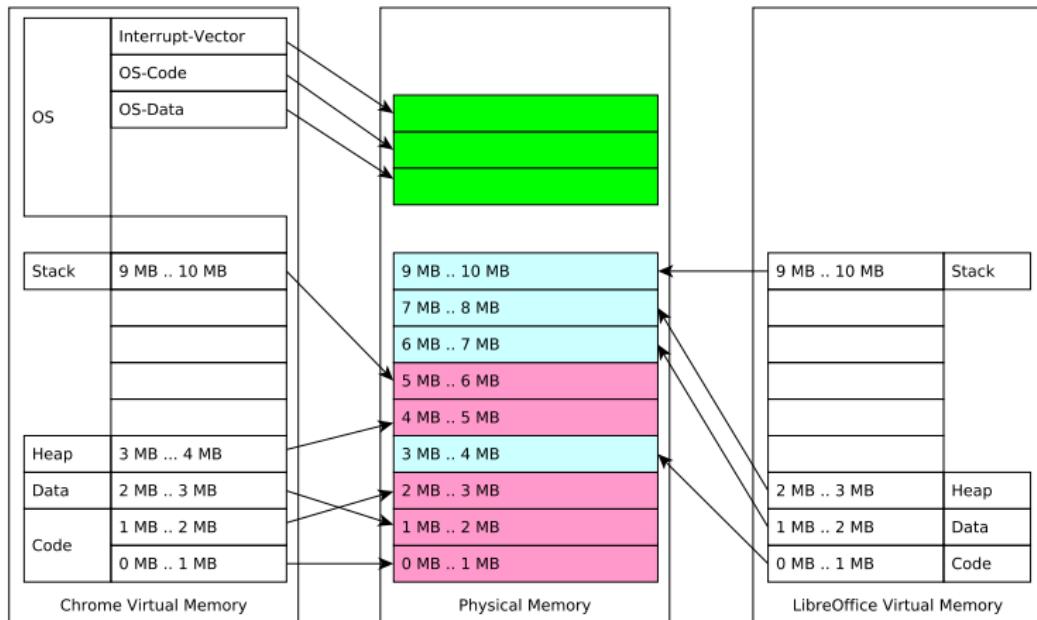
# OS Memory

- ① Question 4: How does the OS access the memory?



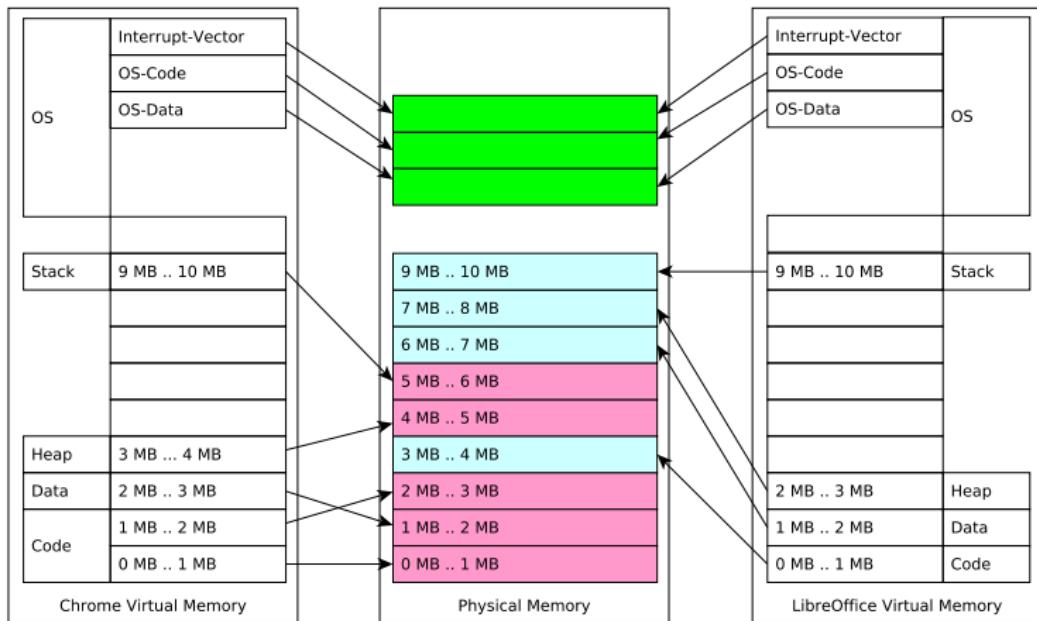
# OS Memory

## ① Question 4: How does the OS access the memory?



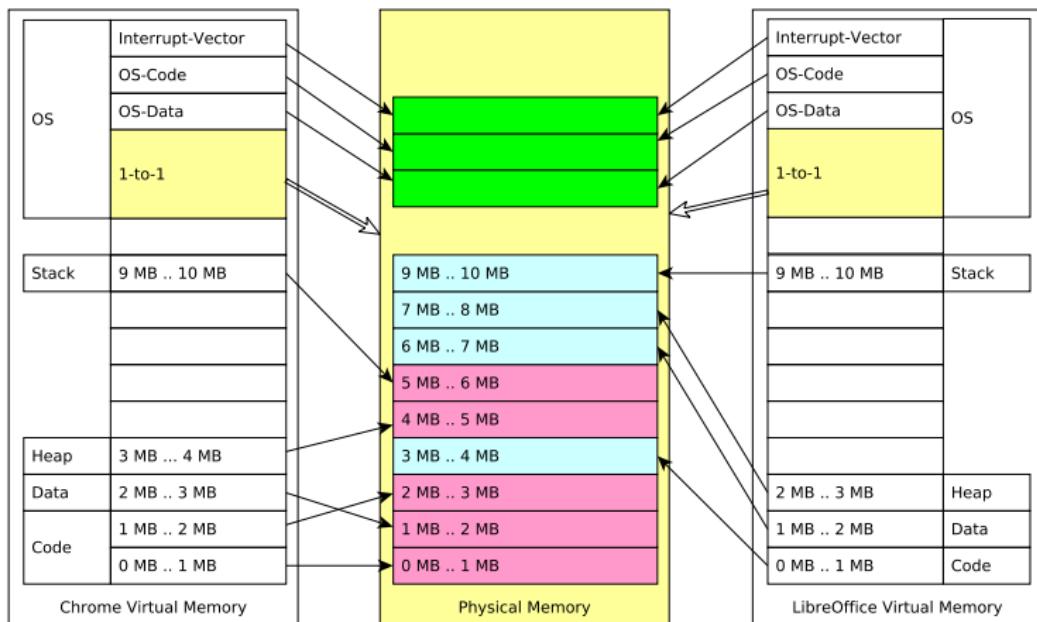
# OS Memory

- ① Question 4: How does the OS access the memory?



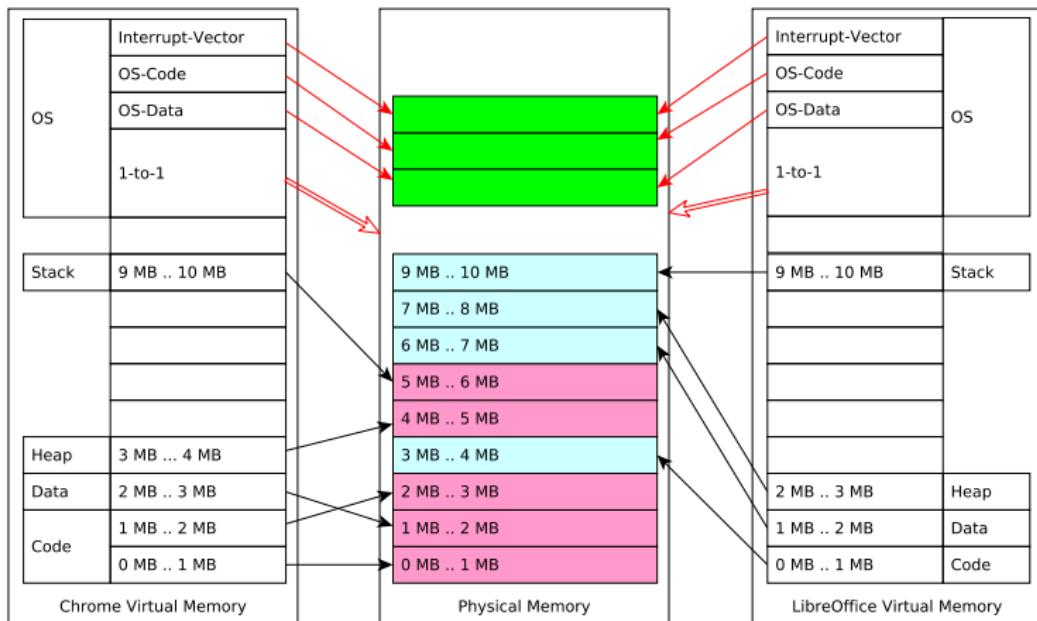
# OS Memory

## ① Question 4: How does the OS access the memory?



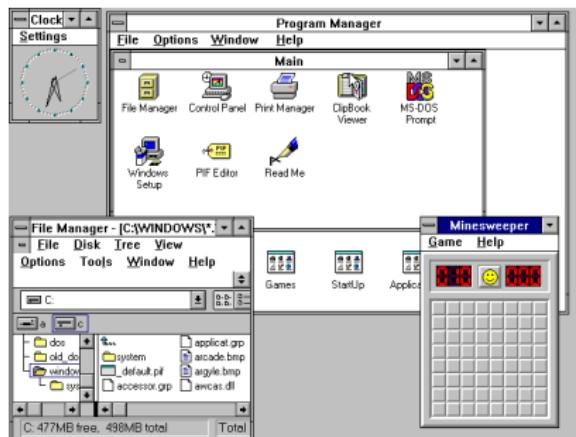
# OS Memory

- ① Question 4: How does the OS access the memory?
- ② MMU enforces access right

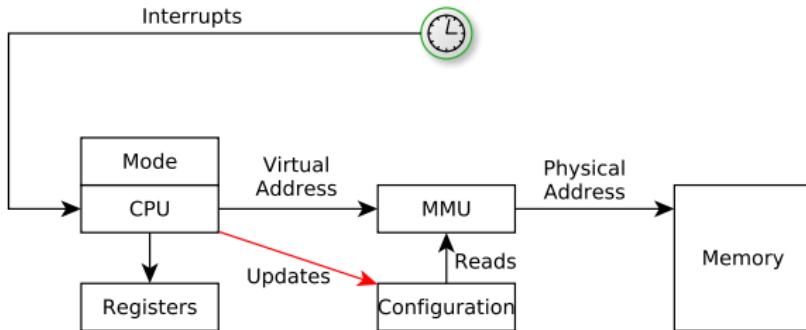


# Non-preemptive multitasking

- Context switch driven by software interrupt
- Cooperative/non-preemptive multitasking

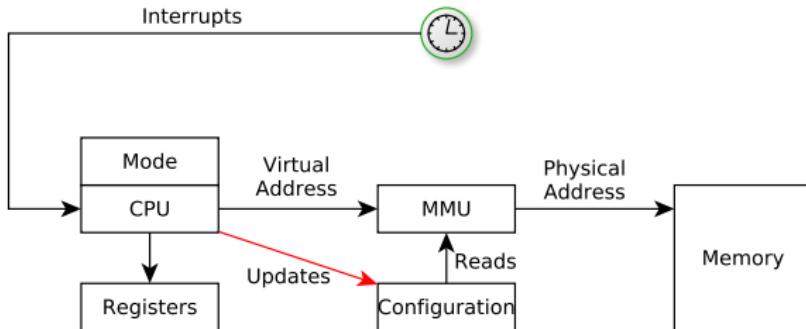


# HW Interrupt



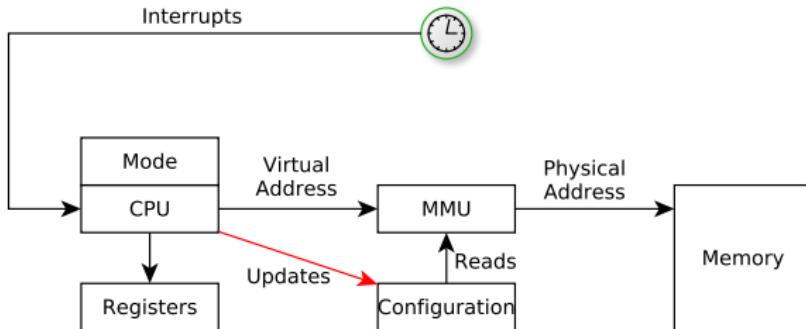
- ① Peripheral (timer) delivers an Hardware Interrupt

# HW Interrupt



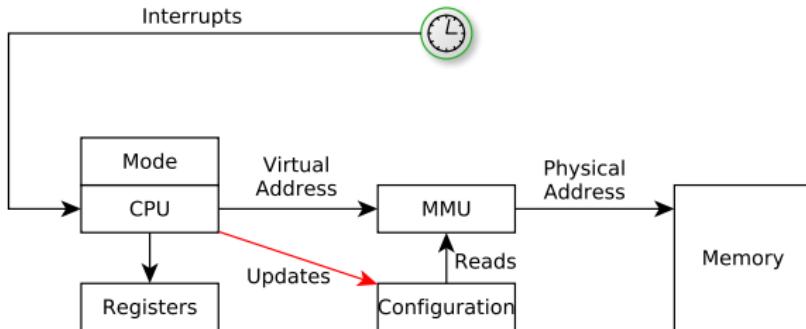
- ① Peripheral (timer) delivers an Hardware Interrupt
- ② HW
  - Mode switches to un-restricted
  - Program counter jumps to the OS “vector table”

# HW Interrupt



- ① Peripheral (timer) delivers an Hardware Interrupt
- ② HW
  - Mode switches to un-restricted
  - Program counter jumps to the OS “vector table”
- ③ OS implements context switch

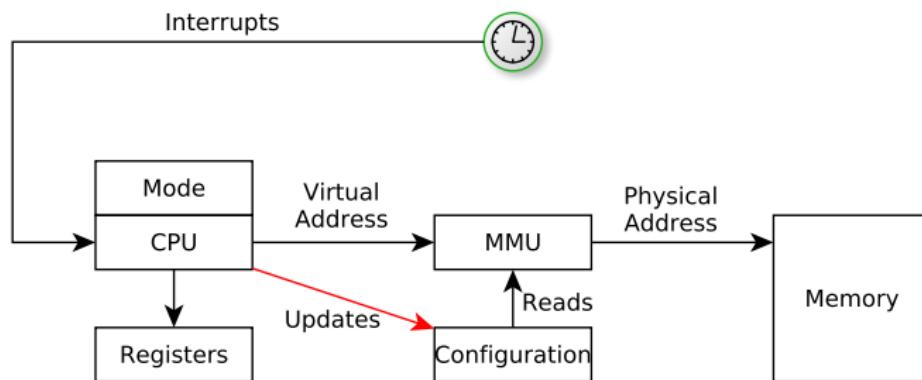
# HW Interrupt



- ① Peripheral (timer) delivers an Hardware Interrupt
- ② HW
  - Mode switches to un-restricted
  - Program counter jumps to the OS “vector table”
- ③ OS implements context switch
- ④ Preemptive multitasking

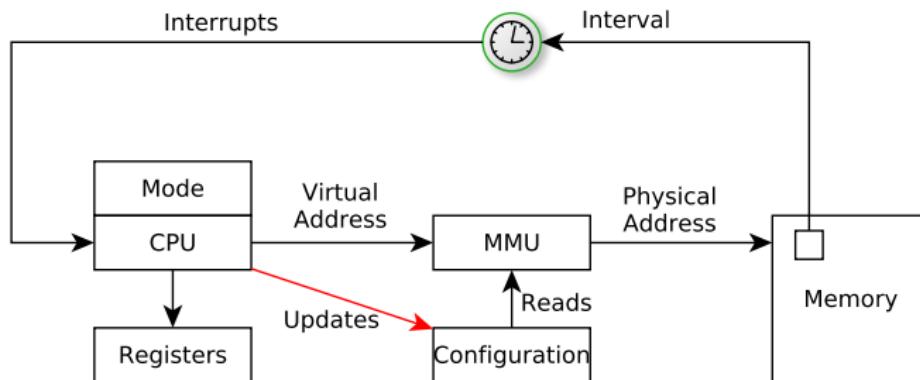
# Peripherals

## ① Interrupts



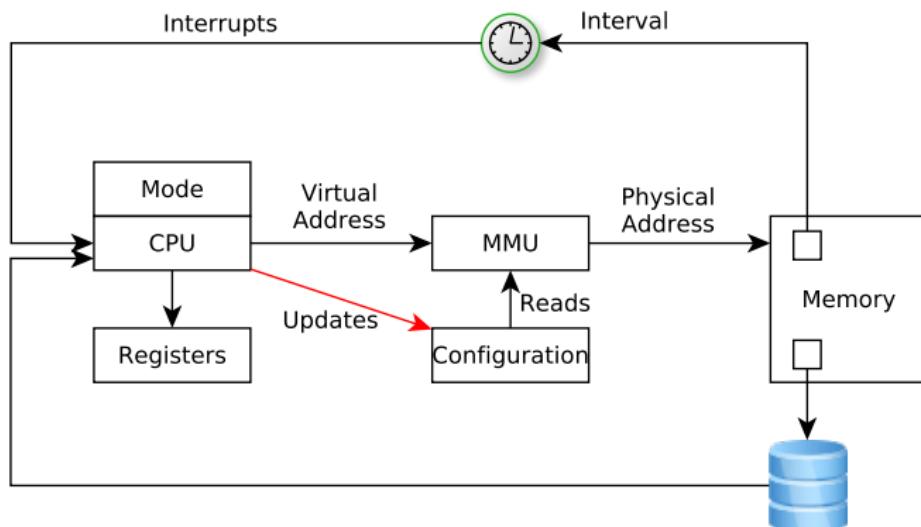
# Peripherals

- ① Interrupts
- ② Memory mapped devices



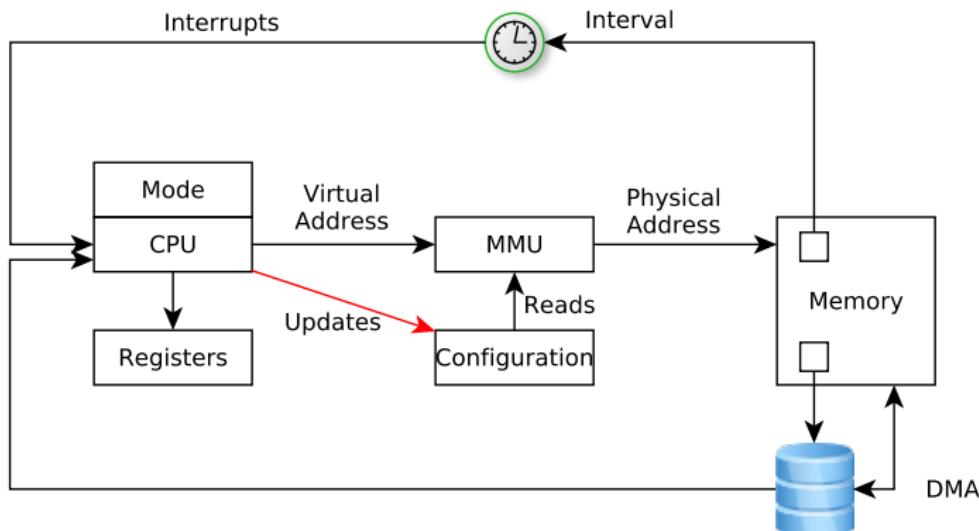
# Peripherals

- ① Interrupts
- ② Memory mapped devices



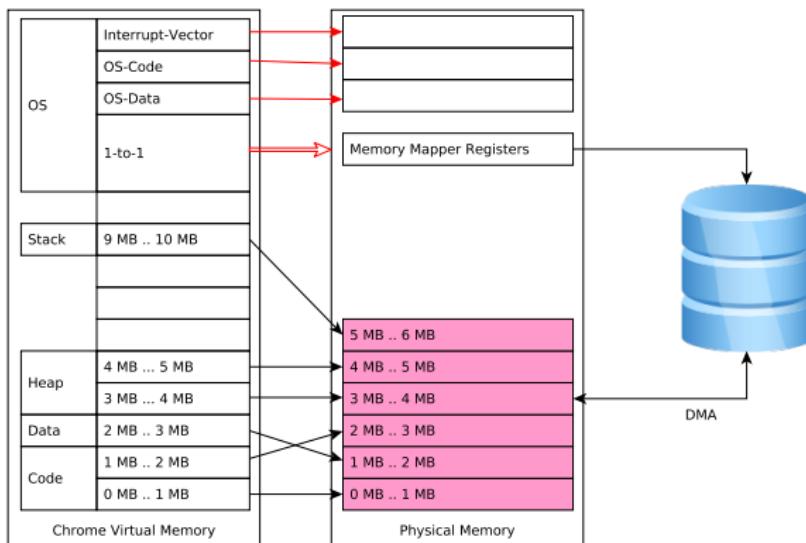
## Peripherals

- ① Interrupts
  - ② Memory mapped devices
  - ③ Direct Memory Access

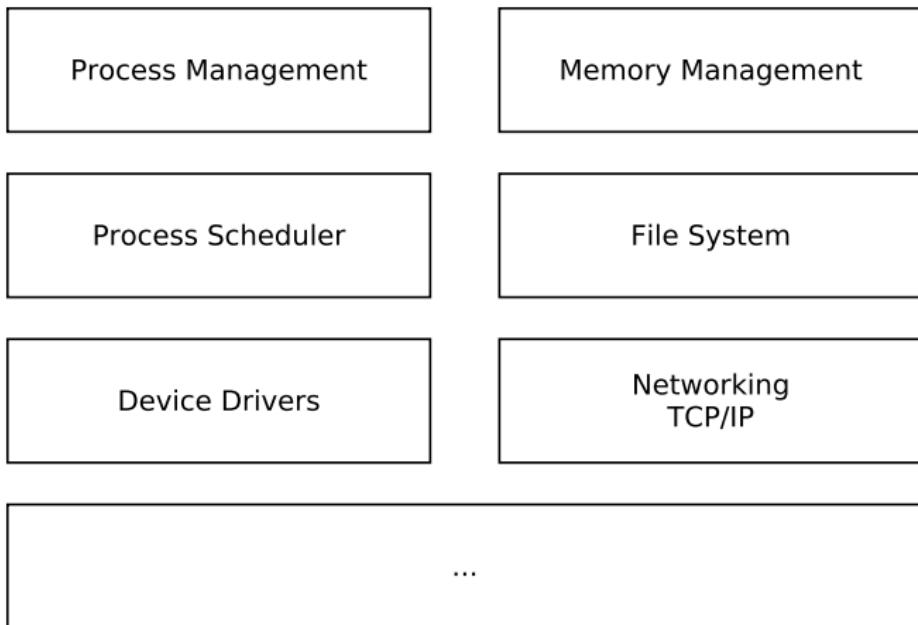


# Peripherals

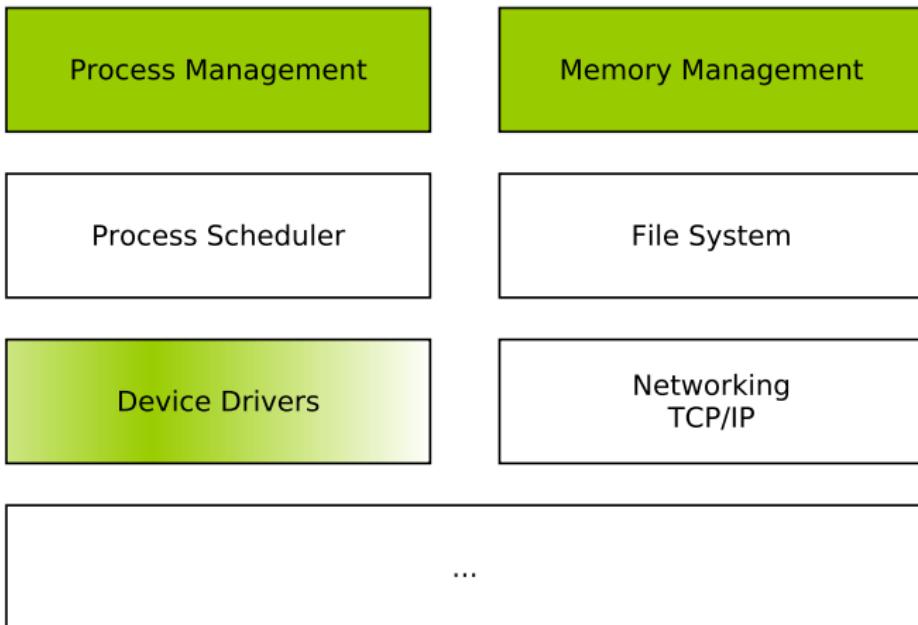
- ① Interrupts
- ② Memory mapped devices
- ③ Direct Memory Access



# OS blocks



# OS blocks



# System calls

- Request OS services
  - File management
  - Device Management
  - Memory allocation
  - Process Control
  - Communication
  - Print to the terminal
- Calling convention
  - Usually invoked by Software Interrupts
  - HW-OS dependent

- Syscall low level (and non-portable)
- Library or API that sits between normal programs and the operating system
- Unix-like: libc, glibc
- Windows NT, that API is part of the Native API, in the ntdll.dll library
- in POSIX: open, read, write, close, wait, exec, fork, exit, mmap

# Summary

- Process memory layout
- Processes vs Thread
- OS is key part of the Trusting Computing Base
- OS allows sharing resources
- OS monitors resources

# Questions?

Questions?